

Parallelizing DIPY model fits with Ray

Ariel Rokem

1 Under construction

2 Abstract

3 Introduction

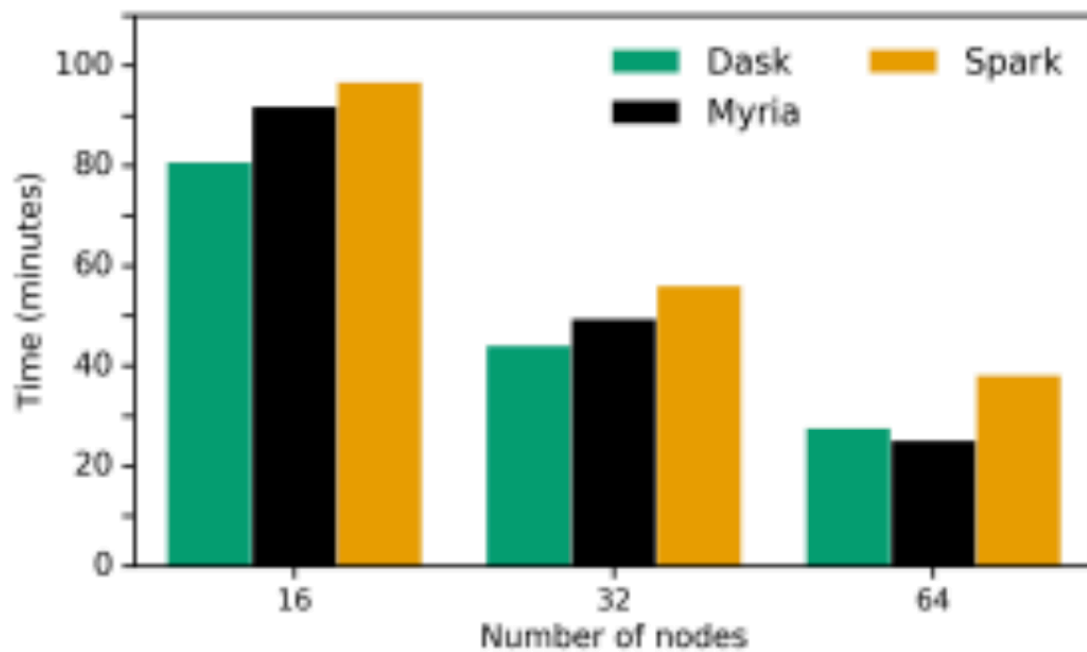
4 Methods

5 Results

Parallelization with `ray` provides improved run-times across a range of different configurations.

Efficiency decreases as a function of number of CPUs, but is still rather high in many configurations.

Ray tends to spill a large amount of data to disk and does not clean up afterwards.



6 Discussion

6.1 Acknowledgments

This work was funded through NIH grant EB027585 (PI: Eleftherios Garyfallidis) and a grant from the Chan Zuckerberg Initiative Essential Open Source Software program (PI: Serge Koudoro).